

Symbolic closed-form. Each channel is fit as an explicit trigonometric formula by ridge regression on D random

Fourier features, $\text{ch}(x, y) = a_0 + \sum_k [a_k \cos(\omega_k \cdot (x, y)) + b_k \sin(\omega_k \cdot (x, y))]$, linear in known analytic basis functions so the picture literally is a readable equation [9]. Fidelity is honest and bounded: a smooth gradient reaches ≈ 57 dB from 512 terms, a hard-edged checkerboard only ≈ 14 , because a finite trigonometric sum cannot render a discontinuity.

Gaussian mixture (2D splatting). The image is a sum of coloured anisotropic Gaussians, $\text{ch}(x, y) = b + \sum_k c_k \exp(-\frac{1}{2} d_k^\top \Sigma_k^{-1} d_k)$, optimized jointly, the image counterpart of the representation behind real-time radiance fields [10], [11].

Thin-plate radial basis. A linear combination of a fixed thin-plate kernel plus an affine plane, $\text{ch}(x, y) = a_0 + a_1 x + a_2 y + \sum_i w_i \phi(\|(x, y) - c_i\|)$ with $\phi(r) = r^2 \log r$, the smoothest interpolant of scattered data [12], solved in closed form.

Chebyshev series. A truncated tensor polynomial series $\sum_{ij} a_{ij} T_j(x) T_i(y)$, $T_k(t) = \cos(k \arccos t)$, fit live in the browser via discretely orthonormalized moments [13]; fidelity saturates on hard edges, the classical moments tradeoff.

Fourier descriptors. The dominant closed contour is traced and transformed into a sum of rotating circles $z(t) = \sum_k c_k e^{i f_k t}$, an exact orderable parametric description of a one-dimensional curve.

Generative latent. The image is a code in the latent space of a variational autoencoder [14]; the useful directions are disentangled and semantic, the second pole of high editability.

III. CROSS-FAMILY BENCHMARK

We bake each representation of six benchmark images (a photograph, a painting, a mathematical-art image, an astronomical image, a texture, and a synthetic gradient) and evaluate two axes (Fig. 1); two further families described above, overcomplete sparse dictionaries and Fourier-descriptor epicycles, appear in the interactive lab but are outside this fixed-budget twelve-family benchmark. Each family is fit at its native working resolution (symbolic, Gabor, Gaussian-mixture and thin-plate at 96, the neural field and primitives at 128, the four transforms and the VAE at 256), so the fixed-budget PSNRs below are not resolution-matched across families and the fidelity comparison is qualitative rather than a strict ranking. *Fixed-budget fidelity*: at a comparable per-unit budget (matched by representational unit, not scalar count: the top 5% of transform coefficients, ≈ 3300 numbers, against ≈ 1200 ellipses, ≈ 2300 neural weights, 250 Gabor atoms, 300 Gaussians, 225 thin-plate kernels, 512 trigonometric terms, or a $4 \times 32 \times 32$ latent), the twelve families land within 26.8–34.8 dB (wavelet 34.8, DCT 33.2, neural field 33.0, Fourier 32.3, Gabor 32.1, primitives 32.0, Gaussian mixture 31.9, thin-plate 29.4, symbolic 29.4, KLT 29.4, Chebyshev 28.9, VAE latent 26.8). Fidelity does not sharply separate the families. *Rate-distortion*: over coefficient fractions from about 2% to 25% the space-localized wavelet dominates (up to 49.4 dB at 25%; DCT leads at the lowest 0.5–1% fractions), then DCT and Fourier, with the patch KLT lower, the classical energy-compaction ordering. *Editability-locality*: we measure how spatially concentrated the effect of a single coefficient

is, operationally the fraction of a single coefficient’s impulse-response energy that falls within a fixed local neighbourhood (1 = perfectly local, 0 = spread across the image; the exact estimator and support size are in the released benchmark code); the global transforms are non-local (Fourier 0.156, DCT 0.228) while the space-localized transforms are local-exact (wavelet 0.999, KLT 1.0). This is the axis on which the families genuinely differ.

IV. THE FINDING: U-SHAPED EDITABILITY

Putting the two axes together yields the finding (Fig. 2). Because fixed-budget fidelity is comparable, what distinguishes the families for a user who wants to *edit* the image is whether a parameter is a stable, meaningful handle, and that is U-shaped across abstraction. Editability is high at the designed-structure pole, geometric primitives (move one ellipse and only that region changes), sparse and Gabor atoms (each a legible local object), and space-localized transforms (wavelet and KLT, locality ≈ 1), because humans built local, meaningful coordinates. It is high again at the learned-manifold pole, the disentangled directions of a generative latent space, because the model learned a low-dimensional manifold of plausible images with semantic axes. And it collapses in between: a single coefficient of a global transform is spatially spread (locality 0.16–0.23, so a local edit is impossible), a raw neural-field weight is entangled (no weight is a named handle), and a constant of a brittle fitted formula is fragile (perturbing it destroys the image, exactly because a finite trigonometric sum has no local support). The organizing statement is that a parameter is both stable and meaningful only when it indexes a low-dimensional manifold of plausible images with locally disentangled coordinates; the designed and the learned poles achieve this by construction and by training, and the entangled middle does not.

V. DISCUSSION AND HONEST SCOPE

The quantitative editability-locality is measured for the four transform families, where the local-versus-global contrast is exact and unambiguous; the U-shape across all twelve families is the organizing thesis, supported by that measurement, by the qualitative locality of each family’s parameters, and by the fixed-budget benchmark that rules out fidelity as the differentiator. The result is descriptive, not a new representation: it says where a representation’s parameters are, and are not, a usable handle, which is the practical question behind interpretable and editable image models. The honest scope is explicit: arbitrary photographs do not reduce to compact, faithful, human-readable equations, the symbolic and formula families make genuine but bounded partial reconstructions (a smooth gradient fits to ≈ 57 dB, a hard edge to ≈ 14), and the benchmark reports those partial results rather than overclaiming. The generative-latent pole is argued from the disentanglement literature and the family’s construction rather than from a locality number, and is labelled as such.

VI. CONCLUSION

We carried one image across twelve mathematical representation families on a common benchmark and found that

Equal-budget fidelity does not distinguish representations; editability does




Fig. 1. The benchmark, read from the committed `_bench/index.json`. Left: peak signal-to-noise ratio at a fixed parameter budget across the twelve families; fidelity is comparable (27–35 dB) and does not distinguish them, coloured by pole (designed-structure local, global or entangled, learned manifold). Right: the editability-locality metric for the four transform families; a single coefficient of a global transform is spatially spread (Fourier 0.16, DCT 0.23) while a coefficient of a space-localized transform is essentially local (wavelet 0.999, KLT 1.0).




Fig. 2. Editability is U-shaped across the abstraction spectrum: high at the designed-structure pole (primitives, sparse atoms, space-localized transforms) and the learned-manifold pole (disentangled generative directions), and low in the entangled middle (global transforms, raw neural-field weights, brittle fitted formulas).

equal-budget fidelity does not distinguish them (27–35 dB) while editability does, and that editability is U-shaped across the abstraction spectrum, high at the designed-structure and learned-manifold poles and low in the entangled middle, with the transform families’ editability-locality measured directly (global 0.16–0.23 versus local-exact ≈ 1). A parameter is a usable handle only when it indexes a low-dimensional manifold with locally disentangled coordinates. All code and committed artifacts are released so the benchmark and the figures are reproducible.

DATA AND CODE AVAILABILITY

Code, the committed benchmark artifacts, and the interactive lab are at https://github.com/fsantibanezleal/CAOS_IMGLAB under the MIT license. Benchmark images

are public-domain or permissively licensed, plus procedurally generated license-clean images; each carries its citation in the repository. Figure 1 is regenerated deterministically from the committed `_bench/index.json` by `manuscripts/image-representation-editability/figure1.py`. The U-shape schematic is hand-authored (`fig-ushape.svg`).

APPENDIX A FIXED-BUDGET FIDELITY TABLE

At the matched budget the peak signal-to-noise ratios are: wavelet 34.8, DCT 33.2, neural field (SIREN) 33.0, Fourier 32.3, Gabor atoms 32.1, primitives 32.0, Gaussian mixture 31.9, thin-plate RBF 29.4, symbolic equation 29.4, KLT (patch) 29.4, Chebyshev series 28.9, VAE latent 26.8 dB, over six benchmark images (photograph, painting, mathematical art, astronomy, texture, synthetic gradient). The full per-image rate-distortion curves and the editability-locality values are committed in `data/derived/_bench/index.json`.

REFERENCES

- [1] N. Ahmed, T. Natarajan, and K. R. Rao, “Discrete cosine transform,” *IEEE Trans. Comput.*, vol. C-23, no. 1, pp. 90–93, 1974, doi:10.1109/T-C.1974.223784.
- [2] S. G. Mallat, “A theory for multiresolution signal decomposition: The wavelet representation,” *IEEE Trans. Pattern Anal. Mach. Intell.*, vol. 11, no. 7, pp. 674–693, 1989, doi:10.1109/34.192463.
- [3] G. K. Wallace, “The JPEG still picture compression standard,” *IEEE Trans. Consum. Electron.*, vol. 38, no. 1, pp. xviii–xxxiv, 1992, doi:10.1109/30.125072.
- [4] S. G. Mallat and Z. Zhang, “Matching pursuits with time-frequency dictionaries,” *IEEE Trans. Signal Process.*, vol. 41, no. 12, pp. 3397–3415, 1993, doi:10.1109/78.258082.

- [5] M. Aharon, M. Elad, and A. Bruckstein, “K-SVD: An algorithm for designing overcomplete dictionaries for sparse representation,” *IEEE Trans. Signal Process.*, vol. 54, no. 11, pp. 4311–4322, 2006, doi:10.1109/TSP.2006.881199.
- [6] J. G. Daugman, “Uncertainty relation for resolution in space, spatial frequency, and orientation optimized by two-dimensional visual cortical filters,” *J. Opt. Soc. Amer. A*, vol. 2, no. 7, pp. 1160–1169, 1985, doi:10.1364/JOSAA.2.001160.
- [7] V. Sitzmann *et al.*, “Implicit neural representations with periodic activation functions,” in *Adv. Neural Inf. Process. Syst. (NeurIPS)*, 2020, arXiv:2006.09661.
- [8] M. Tancik *et al.*, “Fourier features let networks learn high frequency functions in low dimensional domains,” in *Adv. Neural Inf. Process. Syst. (NeurIPS)*, 2020, arXiv:2006.10739.
- [9] K. O. Stanley, “Compositional pattern producing networks: A novel abstraction of development,” *Genet. Program. Evolvable Mach.*, vol. 8, pp. 131–162, 2007, doi:10.1007/s10710-007-9028-8.
- [10] B. Kerbl, G. Kopanas, T. Leimkühler, and G. Drettakis, “3D Gaussian splatting for real-time radiance field rendering,” *ACM Trans. Graph.*, vol. 42, no. 4, 2023, doi:10.1145/3592433.
- [11] X. Zhang *et al.*, “GaussianImage: 1000 FPS image representation and compression by 2D Gaussian splatting,” in *Proc. Eur. Conf. Comput. Vis. (ECCV)*, 2024, arXiv:2403.08551.
- [12] F. L. Bookstein, “Principal warps: Thin-plate splines and the decomposition of deformations,” *IEEE Trans. Pattern Anal. Mach. Intell.*, vol. 11, no. 6, pp. 567–585, 1989, doi:10.1109/34.24792.
- [13] R. Mukundan, S. H. Ong, and P. A. Lee, “Image analysis by Tchebichef moments,” *IEEE Trans. Image Process.*, vol. 10, no. 9, pp. 1357–1364, 2001, doi:10.1109/83.941859.
- [14] D. P. Kingma and M. Welling, “Auto-encoding variational Bayes,” in *Int. Conf. Learn. Represent. (ICLR)*, 2014, arXiv:1312.6114.